

Test

November 13, 2024

```
[35]: # Add the Model Scripts folder to the path
import sys
sys.path.append("Model Scripts")
sys.path.append("Model Weights")
```

```
# Import the necessary libraries
import torch
import random
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt
```

```
[36]: # Import the function to create the model
from Scripts.ArithmeticTransformer import create_arithmetic_transformer
```

```
# Create a model with default parameters
model = create_arithmetic_transformer()
```

```
# Or create a model with custom parameters
model = create_arithmetic_transformer(
    vocab_size=14,
    embed_size=128,
    num_heads=4,
    ff_dim=512,
    num_layers=3,
    max_length=64
)
```

```
[37]: import torch

def load_model(model_path, device=None):
    # Determine device
    if device is None:
        device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

```

# Load checkpoint
print(f"Loading model on {device}")
checkpoint = torch.load(model_path, map_location=device)

# Extract config
config = checkpoint['model_config']
print("Model configuration:", config)

# Create model
model = create_arithmetic_transformer(
    vocab_size=config['vocab_size'],
    embed_size=config['embed_size'],
    num_heads=config['num_heads'],
    ff_dim=config['ff_dim'],
    num_layers=config['num_layers'],
    max_length=config['max_length'],
    dropout=config['dropout']
)

# Load state dict
model.load_state_dict(checkpoint['model_state_dict'])

# Move model to device and set to eval mode
model = model.to(device)
model.eval()

# Get vocab if available
vocab = checkpoint.get('vocab')
inv_vocab = checkpoint.get('inv_vocab')

print(f"Model loaded successfully! Best accuracy: {checkpoint['accuracy']:.4f}")

return model, vocab, inv_vocab, device, config

# Usage:
model_path = './Weights/large_addition_model.pth'
model, vocab, inv_vocab, device, config = load_model(model_path)

```

Loading model on cpu

Model configuration: {'vocab_size': 14, 'embed_size': 512, 'num_heads': 8, 'ff_dim': 2048, 'num_layers': 8, 'max_length': 42, 'dropout': 0.15}

Model loaded successfully! Best accuracy: 1.0000

```

[38]: # If you need to use the model for inference, you'll want these helper
      ↪ functions:
def preprocess_input(input_str, max_length, vocab):

```

```

# Reverse the input string
input_str = input_str[::-1]
# Tokenize
tokens = [vocab[c] for c in input_str if c in vocab]
# Pad
padded = tokens + [vocab['<PAD>']] * (max_length - len(tokens))
return torch.tensor(padded).unsqueeze(0) # Add batch dimension

def decode_output(output_tensor, inv_vocab):
    _, predicted = output_tensor.max(2)
    decoded = []
    for token in predicted[0]:
        token_val = token.item()
        if token_val == vocab['<EOS>']:
            break
        if token_val != vocab['<PAD>']:
            decoded.append(inv_vocab[token_val])
    return ''.join(decoded[::-1]) # Reverse at the end

# Example usage:
def test_addition(num1, num2, model, vocab, inv_vocab, max_length):
    input_str = f"{num1}+{num2}="
    input_tensor = preprocess_input(input_str, max_length, vocab)
    with torch.no_grad():
        output = model(input_tensor)
        result = decode_output(output, inv_vocab)
    print(f"{num1} + {num2} = {result}")
    print(f"Correct result: {num1 + num2}")
    print(f"Model's prediction is {'correct' if int(result) == num1 + num2 else 'incorrect'}")

```

```

[39]: # Test a simple addition
test_addition(123, 456, model, vocab, inv_vocab, config['max_length'])

# or test multiple additions in a loop
test_cases = [
    (5, 7),
    (42, 58),
    (123, 456),
    (1234, 5678)
]

for num1, num2 in test_cases:
    test_addition(num1, num2, model, vocab, inv_vocab, config['max_length'])

```

123 + 456 = 579

Correct result: 579

Model's prediction is correct

```

5 + 7 = 12
Correct result: 12
Model's prediction is correct
42 + 58 = 100
Correct result: 100
Model's prediction is correct
123 + 456 = 579
Correct result: 579
Model's prediction is correct
1234 + 5678 = 6912
Correct result: 6912
Model's prediction is correct

```

```

[40]: # Test a simple addition
test_addition(123, 456, model, vocab, inv_vocab, config['max_length'])

# or test multiple additions in a loop
test_cases = [
    (5, 7),
    (42, 58),
    (123, 456),
    (1234, 5678),
    (10304923, 123123123),
    (123123123, 10304923)
]

for num1, num2 in test_cases:
    test_addition(num1, num2, model, vocab, inv_vocab, config['max_length'])

```

```

123 + 456 = 579
Correct result: 579
Model's prediction is correct
5 + 7 = 12
Correct result: 12
Model's prediction is correct
42 + 58 = 100
Correct result: 100
Model's prediction is correct
123 + 456 = 579
Correct result: 579
Model's prediction is correct
1234 + 5678 = 6912
Correct result: 6912
Model's prediction is correct
10304923 + 123123123 = 133428046
Correct result: 133428046
Model's prediction is correct

```

123123123 + 10304923 = 133428046

Correct result: 133428046

Model's prediction is correct

```
[41]: import torch
import numpy as np
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt
from collections import defaultdict
import seaborn as sns
import random

class ArithmeticModelTester:
    def __init__(self, model, vocab, inv_vocab, max_seq_length):
        self.model = model
        self.vocab = vocab
        self.inv_vocab = inv_vocab
        self.max_seq_length = max_seq_length
        self.model.eval()
        self.results = defaultdict(dict)

    def preprocess_input(self, input_str):
        """Preprocess input string for model"""
        input_str = input_str[::-1] # Reverse string
        tokens = [self.vocab[c] for c in input_str if c in self.vocab]
        padded = tokens + [self.vocab['<PAD>']] * (self.max_seq_length -
        ↪len(tokens))
        return torch.tensor(padded).unsqueeze(0)

    def decode_output(self, output_tensor):
        """Decode model output"""
        _, predicted = output_tensor.max(2)
        decoded = []
        for token in predicted[0]:
            token_val = token.item()
            if token_val == self.vocab['<EOS>']:
                break
            if token_val != self.vocab['<PAD>']:
                decoded.append(self.inv_vocab[token_val])
        return ''.join(decoded[::-1])

    def test_single_digits(self, num_trials=100):
        """Test single digit additions"""
        correct = 0
        for _ in range(num_trials):
            n1 = random.randint(0, 9)
            n2 = random.randint(0, 9)
```

```

        result = self.predict_addition(n1, n2)
        if result == n1 + n2:
            correct += 1
    self.results['single_digits'] = {'accuracy': correct/num_trials}
    return correct/num_trials

def test_commutative_property(self, max_digit=999, num_trials=100):
    """Test if a + b == b + a"""
    correct = 0
    for _ in range(num_trials):
        n1 = random.randint(0, max_digit)
        n2 = random.randint(0, max_digit)
        result1 = self.predict_addition(n1, n2)
        result2 = self.predict_addition(n2, n1)
        if result1 == result2 == (n1 + n2):
            correct += 1
    self.results['commutative'] = {'accuracy': correct/num_trials}
    return correct/num_trials

def test_zero_property(self, max_digit=999, num_trials=100):
    """Test additions with zero"""
    correct = 0
    for _ in range(num_trials):
        n = random.randint(0, max_digit)
        result1 = self.predict_addition(n, 0)
        result2 = self.predict_addition(0, n)
        if result1 == result2 == n:
            correct += 1
    self.results['zero_property'] = {'accuracy': correct/num_trials}
    return correct/num_trials

def test_by_length(self, max_length=5):
    """Test additions with different number lengths"""
    results = {}
    for length in range(1, max_length + 1):
        correct = 0
        trials = 100
        for _ in range(trials):
            n1 = random.randint(10**(length-1), 10**length - 1)
            n2 = random.randint(10**(length-1), 10**length - 1)
            result = self.predict_addition(n1, n2)
            if result == n1 + n2:
                correct += 1
        results[length] = correct/trials
    self.results['length_wise'] = results
    return results

```

```

def test_carries(self):
    """Test additions requiring different numbers of carries"""
    test_cases = [
        (9, 1),      # Single carry
        (99, 1),     # Double carry
        (999, 1),    # Triple carry
        (19, 81),    # Multiple carries
        (999999, 1)  # Many carries
    ]
    results = {}
    for n1, n2 in test_cases:
        result = self.predict_addition(n1, n2)
        results[f"{n1}+{n2}"] = result == n1 + n2
    self.results['carries'] = results
    return results

def predict_addition(self, n1, n2):
    """Make a prediction for n1 + n2"""
    input_str = f"{n1}+{n2}="
    input_tensor = self.preprocess_input(input_str)
    with torch.no_grad():
        output = self.model(input_tensor)
        result_str = self.decode_output(output)
        try:
            return int(result_str)
        except ValueError:
            return None

def visualize_results(self):
    """Visualize test results"""
    plt.figure(figsize=(15, 10))

    # Plot accuracy by number length
    if 'length_wise' in self.results:
        plt.subplot(2, 2, 1)
        lengths = list(self.results['length_wise'].keys())
        accuracies = list(self.results['length_wise'].values())
        plt.plot(lengths, accuracies, marker='o')
        plt.title('Accuracy by Number Length')
        plt.xlabel('Number Length (digits)')
        plt.ylabel('Accuracy')

    # Plot bar chart of different properties
    properties = ['single_digits', 'commutative', 'zero_property']
    accuracies = [self.results[prop]['accuracy'] for prop in properties if
    ↪prop in self.results]
    if accuracies:

```

```

plt.subplot(2, 2, 2)
plt.bar(properties, accuracies)
plt.title('Accuracy by Property')
plt.xticks(rotation=45)
plt.ylabel('Accuracy')

plt.tight_layout()
plt.show()

# Example usage:
"""

# Load your model and necessary components
model = load_model('path_to_model.pth')
vocab = {...} # Your vocabulary
inv_vocab = {...} # Inverse vocabulary
max_seq_length = 42 # Your max sequence length

# Create tester instance
tester = ArithmeticModelTester(model, vocab, inv_vocab, max_seq_length)

# Run tests
tester.test_single_digits()
tester.test_commutative_property()
tester.test_zero_property()
tester.test_by_length()
tester.test_carries()

# Visualize results
tester.visualize_results()
"""

```

```

[41]: "\n# Load your model and necessary components\nmodel =
load_model('path_to_model.pth')\nvocab = {...} # Your vocabulary\ninv_vocab =
{...} # Inverse vocabulary\nmax_seq_length = 42 # Your max sequence
length\n\n# Create tester instance\ntester = ArithmeticModelTester(model, vocab,
inv_vocab, max_seq_length)\n\n# Run tests\ntester.test_single_digits()\ntester.t
est_commutative_property()\ntester.test_zero_property()\ntester.test_by_length()
\ntester.test_carries()\n\n# Visualize results\ntester.visualize_results()\n"

```

```

[45]: import torch
import numpy as np
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt
import seaborn as sns
from tqdm import tqdm
import random
import pandas as pd

```



```

from datetime import datetime

class ModelEvaluator:
    def __init__(self, model, vocab, inv_vocab, max_seq_length, random_seed=42):
        """
        Initialize evaluator with model and set random seed for reproducibility
        """
        self.model = model
        self.vocab = vocab
        self.inv_vocab = inv_vocab
        self.max_seq_length = max_seq_length
        self.model.eval()

        # Set random seeds for reproducibility
        torch.manual_seed(random_seed)
        np.random.seed(random_seed)
        random.seed(random_seed)

        # Store test configuration
        self.test_config = {
            'random_seed': random_seed,
            'timestamp': datetime.now().strftime('%Y-%m-%d_%H-%M-%S'),
            'max_seq_length': max_seq_length
        }

    def generate_test_case(self, min_digits=1, max_digits=20):
        """Generate a random test case with specified digit length range"""
        len1 = random.randint(min_digits, max_digits)
        len2 = random.randint(min_digits, max_digits)

        num1 = random.randint(10**(len1-1), 10**len1 - 1)
        num2 = random.randint(10**(len2-1), 10**len2 - 1)

        return num1, num2, len1, len2

    def preprocess_input(self, input_str):
        """Preprocess input string for model"""
        input_str = input_str[::-1] # Reverse string
        tokens = [self.vocab[c] for c in input_str if c in self.vocab]
        padded = tokens + [self.vocab['<PAD>']] * (self.max_seq_length -
↪len(tokens))
        return torch.tensor(padded).unsqueeze(0)

    def decode_output(self, output_tensor):
        """Decode model output"""
        _, predicted = output_tensor.max(2)
        decoded = []

```

```

for token in predicted[0]:
    token_val = token.item()
    if token_val == self.vocab['<EOS>']:
        break
    if token_val != self.vocab['<PAD>']:
        decoded.append(self.inv_vocab[token_val])
return ''.join(decoded)[::-1]

def evaluate_model(self, num_samples=10000):
    """
    Evaluate model on random test cases and collect detailed metrics
    """
    results = []

    for _ in tqdm(range(num_samples), desc="Evaluating Model"):
        # Generate test case
        num1, num2, len1, len2 = self.generate_test_case()
        true_result = num1 + num2

        # Get model prediction
        input_str = f"{num1}+{num2}="
        input_tensor = self.preprocess_input(input_str)

        with torch.no_grad():
            output = self.model(input_tensor)
            pred_str = self.decode_output(output)

        try:
            pred_result = int(pred_str)
            is_correct = pred_result == true_result
        except ValueError:
            pred_result = None
            is_correct = False

        # Collect detailed information about this test case
        results.append({
            'num1': num1,
            'num2': num2,
            'len1': len1,
            'len2': len2,
            'max_len': max(len1, len2),
            'total_len': len1 + len2,
            'true_result': true_result,
            'predicted_result': pred_result,
            'is_correct': is_correct,
            'carries': self.count_carries(num1, num2)
        })

```

```

    return pd.DataFrame(results)

def count_carries(self, num1, num2):
    """Count the number of carry operations in addition"""
    carry = 0
    carries = 0

    while num1 > 0 or num2 > 0 or carry:
        digit1 = num1 % 10
        digit2 = num2 % 10

        if digit1 + digit2 + carry >= 10:
            carries += 1
            carry = 1
        else:
            carry = 0

        num1 //= 10
        num2 //= 10

    return carries

def analyze_results(self, df):
    """
    Analyze results and generate comprehensive metrics
    """
    metrics = {
        'test_config': self.test_config,
        'overall_accuracy': df['is_correct'].mean(),
        'total_samples': len(df),
        'metrics_by_length': {},
        'metrics_by_carries': {}
    }

    # Analyze by maximum length of operands
    for length in range(1, 21):
        length_df = df[df['max_len'] == length]
        if len(length_df) > 0:
            metrics['metrics_by_length'][length] = {
                'accuracy': length_df['is_correct'].mean(),
                'samples': len(length_df)
            }

    # Analyze by number of carries
    for carries in df['carries'].unique():
        carries_df = df[df['carries'] == carries]

```

```

        metrics['metrics_by_carries'][int(carries)] = {
            'accuracy': carries_df['is_correct'].mean(),
            'samples': len(carries_df)
        }

    return metrics

def visualize_results(self, df, metrics):
    """
    Create visualizations of model performance
    """
    plt.figure(figsize=(20, 10))

    # Plot 1: Accuracy by maximum length
    plt.subplot(2, 2, 1)
    lengths = list(metrics['metrics_by_length'].keys())
    accuracies = [metrics['metrics_by_length'][l]['accuracy'] for l in ↪lengths]
    plt.plot(lengths, accuracies, marker='o')
    plt.title('Accuracy by Maximum Length of Operands')
    plt.xlabel('Maximum Length (digits)')
    plt.ylabel('Accuracy')

    # Plot 2: Accuracy by number of carries
    plt.subplot(2, 2, 2)
    carries = list(metrics['metrics_by_carries'].keys())
    carry_accuracies = [metrics['metrics_by_carries'][c]['accuracy'] for c ↪in carries]
    plt.plot(carries, carry_accuracies, marker='o', color='orange')
    plt.title('Accuracy by Number of Carries')
    plt.xlabel('Number of Carries')
    plt.ylabel('Accuracy')

    # Plot 3: Distribution of test cases
    plt.subplot(2, 2, 3)
    sns.histplot(data=df, x='max_len', bins=20)
    plt.title('Distribution of Test Cases by Maximum Length')
    plt.xlabel('Maximum Length (digits)')

    # Plot 4: Heatmap of length combinations
    plt.subplot(2, 2, 4)
    heatmap_data = pd.crosstab(df['len1'], df['len2'])
    sns.heatmap(heatmap_data, cmap='YlOrRd')
    plt.title('Distribution of Length Combinations')
    plt.xlabel('Length of Second Number')
    plt.ylabel('Length of First Number')

```

```

plt.tight_layout()
plt.show()

# Print summary statistics
print("\nSummary Statistics:")
print(f"Overall Accuracy: {metrics['overall_accuracy']:.4f}")
print(f"Total Samples: {metrics['total_samples']}")

```

```

[ ]: import torch
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt
import seaborn as sns
from tqdm import tqdm
import random
import pandas as pd
from datetime import datetime

# Recreate the model architecture (you'll need to import your model class)
from Scripts.ArithmeticTransformer import *

# First, let's load the model correctly
def load_model_and_config(model_path):
    """Load the model and its configuration"""
    # Load checkpoint
    print(f>Loading model on {device}")
    checkpoint = torch.load(model_path, map_location=device)

    # Extract config
    config = checkpoint['model_config']
    print("Model configuration:", config)

    # Create model
    model = create_arithmetic_transformer(
        vocab_size=config['vocab_size'],
        embed_size=config['embed_size'],
        num_heads=config['num_heads'],
        ff_dim=config['ff_dim'],
        num_layers=config['num_layers'],
        max_length=config['max_length'],
        dropout=config['dropout']
    )

    # Load state dict
    model.load_state_dict(checkpoint['model_state_dict'])

    # Move model to device and set to eval mode

```

```

model = model.to(device)
model.eval()

# Get vocab if available
vocab = checkpoint.get('vocab')
inv_vocab = checkpoint.get('inv_vocab')

print(f"Model loaded successfully! Best accuracy: {checkpoint['accuracy']:.4f}")

return model, vocab, inv_vocab, config['max_length']

# [Previous ModelEvaluator class code remains the same]

```

```

[52]: # SMALL MODEL EVALUATION
# Load model and components
model_path = './Weights/small_addition_model.pth' # Update with your model path
model, vocab, inv_vocab, max_seq_length = load_model_and_config(model_path)

# Create evaluator with specific random seed
evaluator = ModelEvaluator(model, vocab, inv_vocab, max_seq_length,
    random_seed=42)

# Run evaluation
results_df = evaluator.evaluate_model(num_samples=10000)
metrics = evaluator.analyze_results(results_df)

# Visualize results
evaluator.visualize_results(results_df, metrics)

# Save results
timestamp = datetime.now().strftime('%Y-%m-%d_%H-%M-%S')
results_df.to_csv(f'model_evaluation_{timestamp}.csv')

# Print some interesting statistics
print("\nDetailed Statistics:")
print(f"Number of correct predictions: {results_df['is_correct'].sum()}")
print(f"Average number of carries: {results_df['carries'].mean():.2f}")

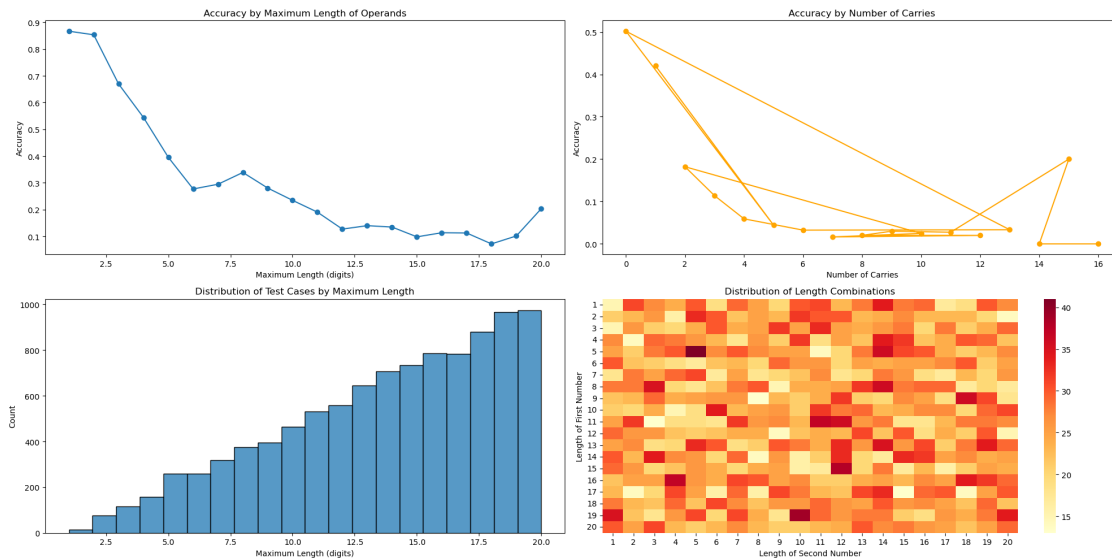
# Print performance by length ranges
print("\nPerformance by length ranges:")
length_ranges = [(1,5), (6,10), (11,15), (16,20)]
for start, end in length_ranges:
    mask = (results_df['max_len'] >= start) & (results_df['max_len'] <= end)
    acc = results_df[mask]['is_correct'].mean()
    print(f"{start}-{end} digits: {acc:.4f}")

```

Loading model on cpu

Model configuration: {'vocab_size': 14, 'embed_size': 64, 'num_heads': 2, 'ff_dim': 256, 'num_layers': 2, 'max_length': 42, 'dropout': 0.1}
 Model loaded successfully! Best accuracy: 0.7937

Evaluating Model: 100%| | 10000/10000 [00:44<00:00, 222.32it/s]



Summary Statistics:

Overall Accuracy: 0.1820
 Total Samples: 10000

Detailed Statistics:

Number of correct predictions: 1820
 Average number of carries: 3.65

Performance by length ranges:

1-5 digits: 0.5507
 6-10 digits: 0.2826
 11-15 digits: 0.1351
 16-20 digits: 0.1221

```
[51]: # MEDUIM MODEL EVALUATION
# Load model and components
model_path = './Weights/medium_addition_model.pth' # Update with your model_
↳ path
model, vocab, inv_vocab, max_seq_length = load_model_and_config(model_path)

# Create evaluator with specific random seed
```

```

evaluator = ModelEvaluator(model, vocab, inv_vocab, max_seq_length,
    ↪random_seed=42)

# Run evaluation
results_df = evaluator.evaluate_model(num_samples=10000)
metrics = evaluator.analyze_results(results_df)

# Visualize results
evaluator.visualize_results(results_df, metrics)

# Save results
timestamp = datetime.now().strftime('%Y-%m-%d_%H-%M-%S')
results_df.to_csv(f'model_evaluation_{timestamp}.csv')

# Print some interesting statistics
print("\nDetailed Statistics:")
print(f"Number of correct predictions: {results_df['is_correct'].sum()}")
print(f"Average number of carries: {results_df['carries'].mean():.2f}")

# Print performance by length ranges
print("\nPerformance by length ranges:")
length_ranges = [(1,5), (6,10), (11,15), (16,20)]
for start, end in length_ranges:
    mask = (results_df['max_len'] >= start) & (results_df['max_len'] <= end)
    acc = results_df[mask]['is_correct'].mean()
    print(f"{start}-{end} digits: {acc:.4f}")

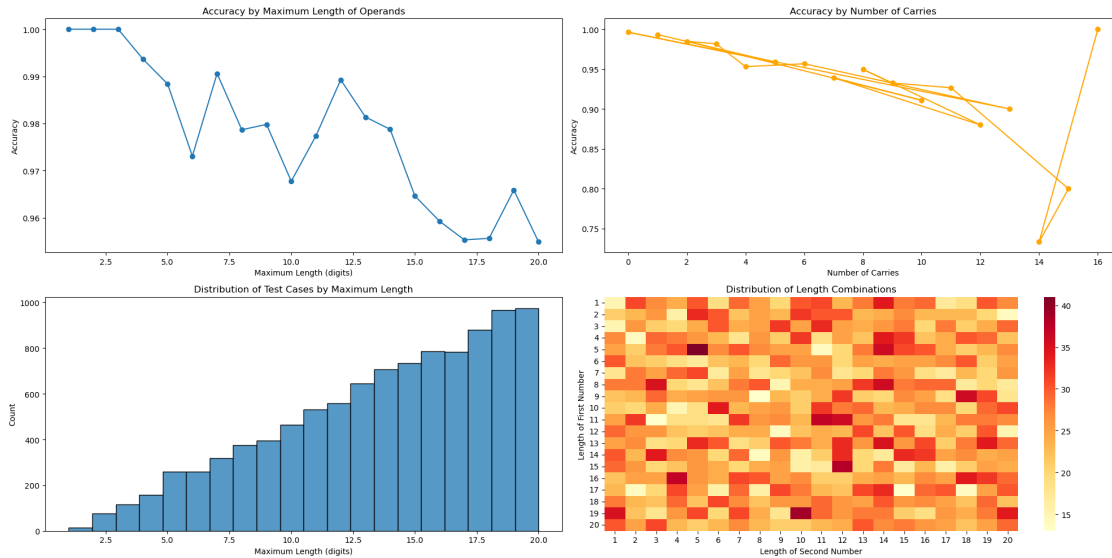
```

Loading model on cpu

Model configuration: {'vocab_size': 14, 'embed_size': 256, 'num_heads': 4, 'ff_dim': 1024, 'num_layers': 4, 'max_length': 42, 'dropout': 0.1}

Model loaded successfully! Best accuracy: 0.9987

Evaluating Model: 100%| | 10000/10000 [01:44<00:00, 96.03it/s]



Summary Statistics:

Overall Accuracy: 0.9701

Total Samples: 10000

Detailed Statistics:

Number of correct predictions: 9701

Average number of carries: 3.65

Performance by length ranges:

1-5 digits: 0.9936

6-10 digits: 0.9774

11-15 digits: 0.9776

16-20 digits: 0.9583

```
[53]: # LARGE MODEL EVALUATION
# Load model and components
model_path = './Weights/large_addition_model.pth' # Update with your model path
model, vocab, inv_vocab, max_seq_length = load_model_and_config(model_path)

# Create evaluator with specific random seed
evaluator = ModelEvaluator(model, vocab, inv_vocab, max_seq_length,
    random_seed=42)

# Run evaluation
results_df = evaluator.evaluate_model(num_samples=10000)
metrics = evaluator.analyze_results(results_df)

# Visualize results
```

```

evaluator.visualize_results(results_df, metrics)

# Save results
timestamp = datetime.now().strftime('%Y-%m-%d_%H-%M-%S')
results_df.to_csv(f'model_evaluation_{timestamp}.csv')

# Print some interesting statistics
print("\nDetailed Statistics:")
print(f"Number of correct predictions: {results_df['is_correct'].sum()}")
print(f"Average number of carries: {results_df['carries'].mean():.2f}")

# Print performance by length ranges
print("\nPerformance by length ranges:")
length_ranges = [(1,5), (6,10), (11,15), (16,20)]
for start, end in length_ranges:
    mask = (results_df['max_len'] >= start) & (results_df['max_len'] <= end)
    acc = results_df[mask]['is_correct'].mean()
    print(f"{start}-{end} digits: {acc:.4f}")

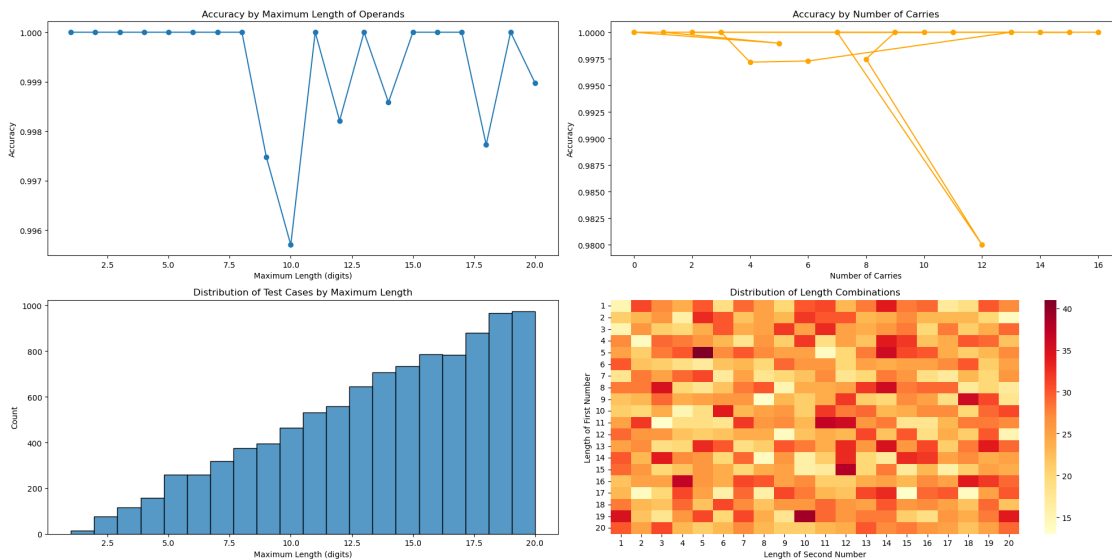
```

Loading model on cpu

Model configuration: {'vocab_size': 14, 'embed_size': 512, 'num_heads': 8, 'ff_dim': 2048, 'num_layers': 8, 'max_length': 42, 'dropout': 0.15}

Model loaded successfully! Best accuracy: 1.0000

Evaluating Model: 100%| | 10000/10000 [13:21<00:00, 12.47it/s]



Summary Statistics:

Overall Accuracy: 0.9992

Total Samples: 10000

Detailed Statistics:

Number of correct predictions: 9992

Average number of carries: 3.65

Performance by length ranges:

1-5 digits: 1.0000

6-10 digits: 0.9983

11-15 digits: 0.9994

16-20 digits: 0.9993

[]: